

Benjamin Feuer

Email: bfeuer@stanford.edu • Website: <https://penfever.github.io>

EDUCATION

Stanford University, Palo Alto, CA — Sept. 2025 – Jun. 2026 Postdoctoral Scholar in Computer Science. Advised by Dr. Ludwig Schmidt.

New York University, New York, NY — Sept. 2021 – May 2025 M.S. / Ph.D. in Computer Science. Advised by Dr. Chinmay Hegde.

Columbia University, New York, NY — Sept. 2008 – Dec. 2012 M.F.A. in Film Production.

Wesleyan University, Middletown, CT — Sept. 2002 – May 2006 B.A. in Film History.

SELECTED PUBLICATIONS

- E Guha, R Marten, ..., B Feuer, ..., A Dimakis, L Schmidt. [Open Thoughts: Data Recipes for Reasoning Models](#). *ICLR (Oral)*, 2026.
- B Feuer, C Hegde. [WildChat-50M: A Deep Dive into the Role of Synthetic Data in Post-Training](#). *ICML*, 2025.
- B Feuer, M Goldblum, T Datta, S Nambiar, R Besaleli, S Dooley, M Cembalest, J P Dickerson. [Style Outweighs Substance: Failure Modes of LLM Judges in Alignment Benchmarking](#). *ICLR*, 2025.
- C White, S Dooley, M Roberts, A Pal, B Feuer, S Jain, R Shwartz-Ziv, N Jain, K Saifullah, S Naidu, C Hegde, Y LeCun, T Goldstein, W Neiswanger, M Goldblum. [LiveBench: A Challenging, Contamination-Free LLM Benchmark](#). *ICLR (Spotlight)*, 2025. In *OpenAI o3-mini T. Report*, *Phi-4 T. Report*, *Gemini 2.0 T. Report*, *Qwen 3 T. Report*, etc.
- B Feuer, C Yang, Z Jubery, Z K. Deng, A Nakkab, M Z Hasan, S Chiranjeevi, K O Marshall, N Baishnab, A K Singh, A Singh, S Sarkar, N Merchant, C Hegde, B Ganapathysubramanian. [BioTrove: A Large Multimodal Dataset Enabling AI for Biodiversity](#). *NeurIPS (Spotlight)*, 2024.

RECENT PUBLICATIONS

- E Guha, R Marten, ..., A Dimakis*, B Feuer*, L Schmidt* (co-last authors). [Open-Thoughts Agent](#). *Blog post*, 2026.
- B Feuer, L Purucker, O Elachqar, C Hegde. [MARVIS: Modality Adaptive Reasoning over VISualizations](#). *arXiv*, 2025.
- J Freire, G Fan, B Feuer, ..., E Wu (alphabetical author order). [Large Language Models for Data Discovery and Integration: Challenges and Opportunities](#). *IEEE Data Eng. Bull. Vol 49*, pp. 3–31, 2025.
- B Feuer, C Tseng, A Lathe, O Elachqar, J P Dickerson. [When Judgment Becomes Noise: How Design Failures in LLM Judge Benchmarks Silently Undermine Validity](#). *arXiv*, 2025.
- K Arabi, B Feuer, R T Witter, C Hegde, N Cohen. [Hidden in the Noise: Two-Stage Robust Watermarking for Images](#). *ICLR*, 2025.
- B Feuer, J Xu, N Cohen, P Yubeaton, G Mittal, C Hegde. [SELECT: Large-Scale Benchmark of Data Curation Strategies for Image Recognition](#). *NeurIPS*, 2024.
- B Feuer, Y Liu, C Hegde, J Freire. [ArcheType: A Novel Framework for Open-Source Column Type Annotation using Large Language Models](#). *VLDB*, 2024.
- B Feuer, R T Schirrmeister, V Cherepanova, C Hegde, F Hutter, M Goldblum, N Cohen, C White. [TuneTables: Context Optimization for Scalable Prior-Data Fitted Networks](#). *NeurIPS*, 2024.
- D McElfresh, S Khandagale, J Valverde, V Prasad C, B Feuer, C Hegde, G Ramakrishnan, M Goldblum, C White. [When Do Neural Nets Outperform Boosted Trees on Tabular Data?](#) *NeurIPS*, 2023.
- B Feuer, A Joshi, M Pham, C Hegde. [Distributionally Robust Classification on a Data Budget](#). *Transactions on Machine Learning Research*, 2023.
- B Feuer, A Joshi, M Cho, K Jani, S Chiranjeevi, ZK Deng, A Balu, A K Singh, S Sarkar, N Merchant, A Singh, B Ganapathysubramanian, C Hegde. [Zero-Shot Insect Detection via Weak Language Supervision](#). *The Plant Phenome Journal (2024)*; *AAAI Workshop on AI for Agriculture and Food Sciences (AIAFS)*, 2023 (Oral).

AWARDS, HONORS, INVITED TALKS

- **Outstanding Dissertation Award, Science and Technology** — NYU, 2026
- **Towards Provably Trustworthy LLM Judges via Bias-Bounded Evaluation** — EFML, Stanford, 2026

- **OpenThoughts-Agent** — Laude Institute at NeurIPS, 2025
- **Annus Mirabilis: A Year of Growth in LLM Research** — David Bau Lab, Northeastern University; JetBrains; Tom Goldstein Lab, UMD. 2025
- **Pearl Brownstein Award (Best Doctoral Research)** — New York University. 2025
- **Style and Substance: Failure Modes of LLM Judges** — NLP, Penn State University. 2024
- **Mitigating Biodiversity Loss with Frontier AI** — Machine Learning, Rensselaer Polytechnic Institute. 2024
- **Top Reviewer Award** — NeurIPS. 2024
- **Style and Substance: Failure Modes of LLM Judges** — AI In Research Sem., New York University. 2024
- **Deborah M. Rosenthal Award (Best CS Qualifying Exam)** — NYU. 2024
- **Big LLMs for Big Data** — Big Data, NYU. 2024

TEACHING

Recitation Leader, New York University

- *Number Theory* — Spring 2022
- *Discrete Mathematics* — Fall 2021, Spring 2022

SERVICE

- **Workshop Lead Organizer:** [Dataworld @ ICML 2025](#); [DCVLR @ NeurIPS 2025](#).
- **Area Chair:** ICML 2026, NeurIPS 2026.
- **Reviewer:** ICML, ICLR, NeurIPS, TMLR — 2024, 2025, 2026. Top Reviewer Award, NeurIPS 2024.
- **Open Source Contributor:** [Llama Factory](#) (60k GitHub stars); [Oumi](#) (8.5k stars); [Harbor](#) (1.5k stars); [SkyRL](#) (2k stars).

PROFESSIONAL EXPERIENCE

OpenAthena, New York, NY (hybrid) — Member of Technical Staff, Jun. 2026 – Present. Leading post-training efforts on the Marin project (fully open-source foundation models up to 720B parameters) in collaboration with Percy Liang at Stanford University.

Oumi.AI, Seattle, WA (remote) — Research Fellow, Jan. 2025 – Present Series A startup; building the “Linux of artificial intelligence”; #1 trending GitHub repo, 8,500+ stars. Conducted original research and developed code for model post-training.

State of Place, Boston, MA (remote) — Research Consultant, Sept. 2024 – Apr. 2025 Led technical team on a Phase 1 SBIR grant with the US DOT. Developed novel, scalable computer vision algorithms for built-environment features and smart cities.

Arthur.AI, New York, NY (hybrid) — Research Fellow, May – Aug. 2024 Series B startup; AI observability platform, including the world’s first LLM firewall. Conducted original research into failure modes of LLM judges for alignment benchmarking.